

Enterprise Artificial Intelligence & Data Policy 2026

Document Version: 3.4 | Effective Date: January 1, 2026 | Classification: Internal Confidential

1. Executive Summary & Purpose

This policy document establishes mandatory guidelines for the usage of Generative AI, Retrieval-Augmented Generation (RAG) systems, Large Language Models (LLMs), and automated data processing pipelines across all enterprise operations. All employees, contractors, and third-party vendors must comply with these protocols to ensure data privacy, intellectual property protection, and regulatory compliance with AI governance frameworks.

2. Approved AI Models & API Usage

2.1 Commercial API Providers: Only enterprise-tier agreements with OpenAI (GPT-4o, text-embedding-3-small), Anthropic (Claude 3.5 Sonnet), and Google Cloud Vertex AI are authorized for processing internal corporate data. Consumer-grade free web interfaces (such as public ChatGPT or Claude web chat) are strictly prohibited for uploading proprietary source code, customer PII, or financial statements.

2.2 Self-Hosted & Local Models: Open-source models deployed via HuggingFace Transformers (e.g., Llama 3, Mistral 7B, Sentence-Transformers all-MiniLM-L6-v2) may be hosted on internal private cloud instances provided vector storage is encrypted at rest using AES-256.

3. Data Classification & Security Levels

Data ingested into RAG vector stores must be tagged according to the following 4 security levels:

Level	Classification	RAG Vector Store Permission	Allowed Embedding Models
L1	Public	Unrestricted internal & public endpoints	HuggingFace, OpenAI, Any
L2	Internal Business	All authenticated employee instances	Enterprise OpenAI, Local HuggingFace
L3	Confidential PII	VPC Isolated Vector Database only	Self-Hosted Local Embeddings Only
L4	Strictly Secret / Trade Secret	PROHIBITED from LLM/RAG Indexing	None (No AI Indexing Permitted)

4. RAG Vector Database & Document Storage Rules

4.1 Retention Period: Document embeddings stored in vector databases (Chroma, Supabase pgvector, Pinecone, or FAISS) must undergo an automated cleanup cycle every 90 days unless tagged for permanent archival retention.

4.2 Chunking Standards: To optimize retrieval accuracy and eliminate hallucination risks, document splitters must adhere to a maximum chunk size of 1,000 characters with an overlap of 150 characters. Chunk metadata MUST preserve document title, section heading, author, and exact page number.

4.3 Deduplication & Hashing: Every incoming PDF must be verified using SHA-256 checksum hashing prior to vector embedding to prevent duplicate index clutter and minimize embedding API costs.

5. Employee Responsibilities & Verification

5.1 Human-in-the-Loop Review: All external communications, code generated by LLMs, and client-facing summaries produced by RAG chatbots must be reviewed and verified by a qualified human team member before publishing or execution.

5.2 Incident Reporting: Any suspected data leak, accidental upload of Level 4 Restricted data to an AI model, or hallucination leading to operational error must be reported to security@enterprise-ai.com within 2 hours of discovery.

6. Compliance Penalties

Non-compliance with this policy may result in revocation of AI tool credentials, formal disciplinary action up to termination of employment, and legal liability under applicable data protection laws (GDPR, CCPA, and EU AI Act 2024).